

## CS89/189: Deep Learning Generalization and Robustness - Winter 2024 Homework#1

Yenkai Huang

## Explain why sometimes the margin can be negative.

The margin in the context of machine learning classification can indeed be negative. The margin, defined as  $\mu(f(x), y) = f(x)[y] - \max_{\{i \neq y\}} f(x)[i]$ , measures the difference between the correct class's prediction score and the highest prediction score of the incorrect classes. A negative margin occurs when the highest score among the incorrect classes exceeds the score of the correct class. This indicates that the model has a higher confidence in an incorrect classification than the true class, which reflects on the model's current inability to correctly distinguish between classes for that particular instance.

## Your proof of the CNN output volume formula

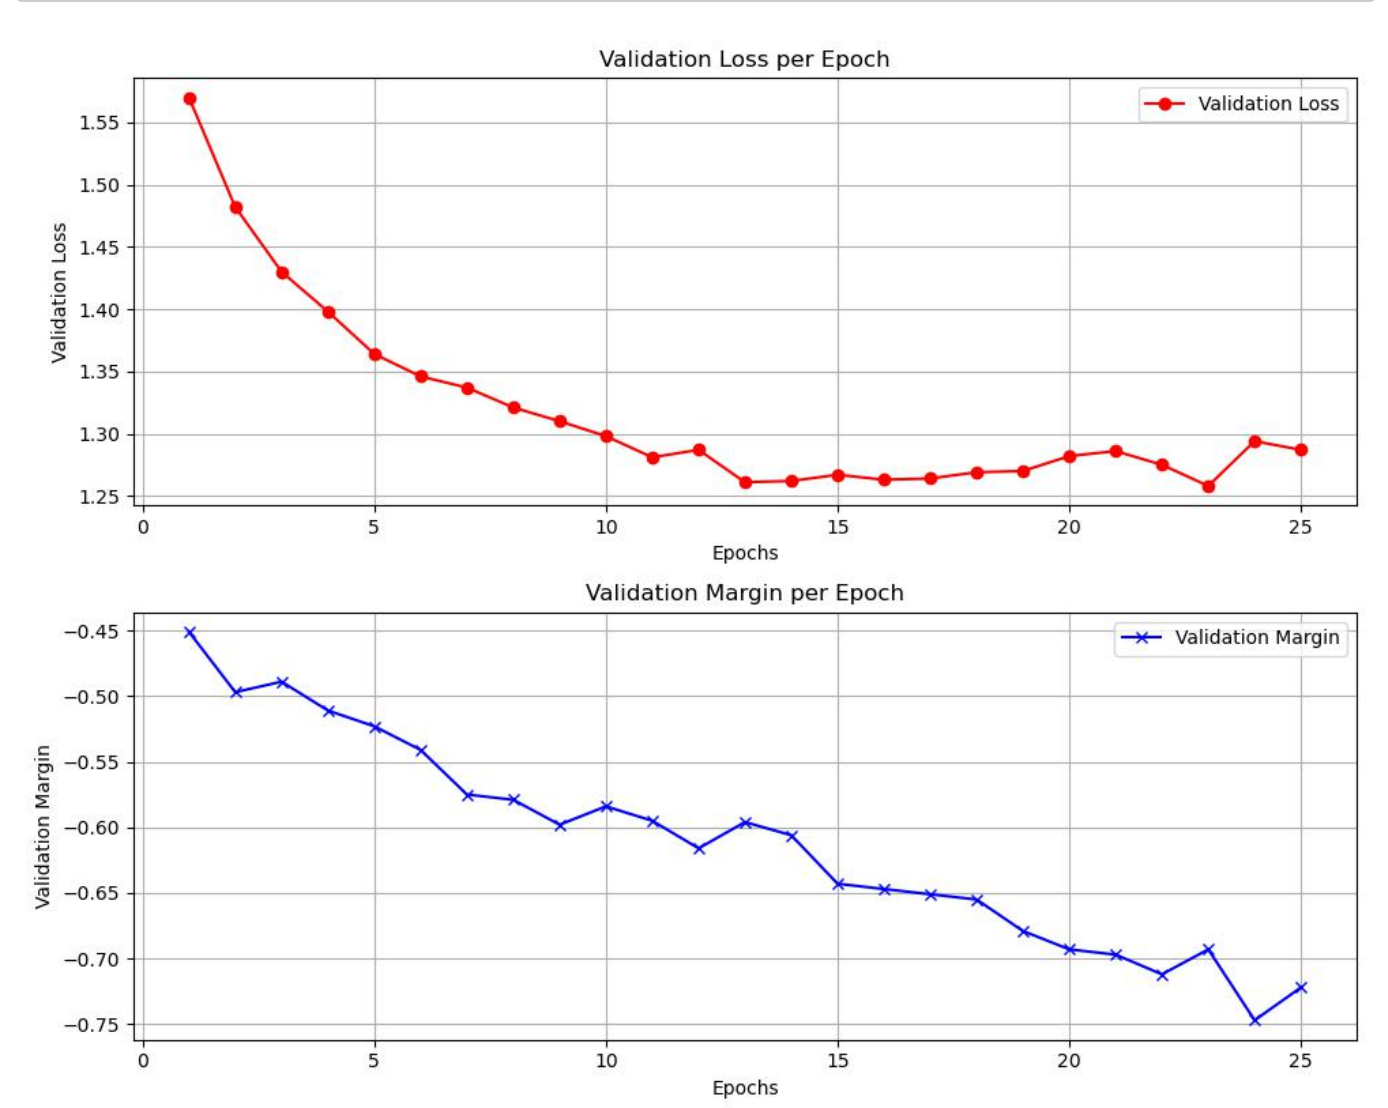
1. When you apply  $n \times n$  filter on a  $N \times N$  input, the output will be  $N - n + 1$
2. If you add the  $Z$  padding (on all sides of edges) the input to expand the size of the output, the output will be  $N - n + 2 \times Z + 1$
3. If we add  $S$  stride to the filtering process, the output will be  $(N - n + 2 \times Z) / S + 1$
4. Having the same filtering process with  $V'$  filters, we have the final output of  $M \times M \times V'$  where  $M = (N - n + 2 \times Z) / S + 1$

## Report your

### Model 1

```
===== Model 1 =====
Training loss: 0.820   Training margin -0.505   Training accuracy: 0.726
Validation accuracy: 0.564

Frobenius1:      20.7
Frobenius2:      6.73
Distance1:       8.99
Distance2:       5.45
Spectral1:       1.91
Spectral2:       3.14
Fro_Fro:         1.4e+02
L1max_L1max:     2.28e+03
Spec_Dist:       32.9
Dist_Spec:       9.02e+02
Spec_Dist_sum:   9.35e+02
Spec_L1max:      32.7
L1max_Spec:      8.79e+02
Spec_L1max_sum:  9.12e+02
Dist_Fro:        60.5
#parameter:      3.16e+06
VC bound:        9.19e+09
L1max bound:     4.59e+11
Your bound:      1.87e+11
```

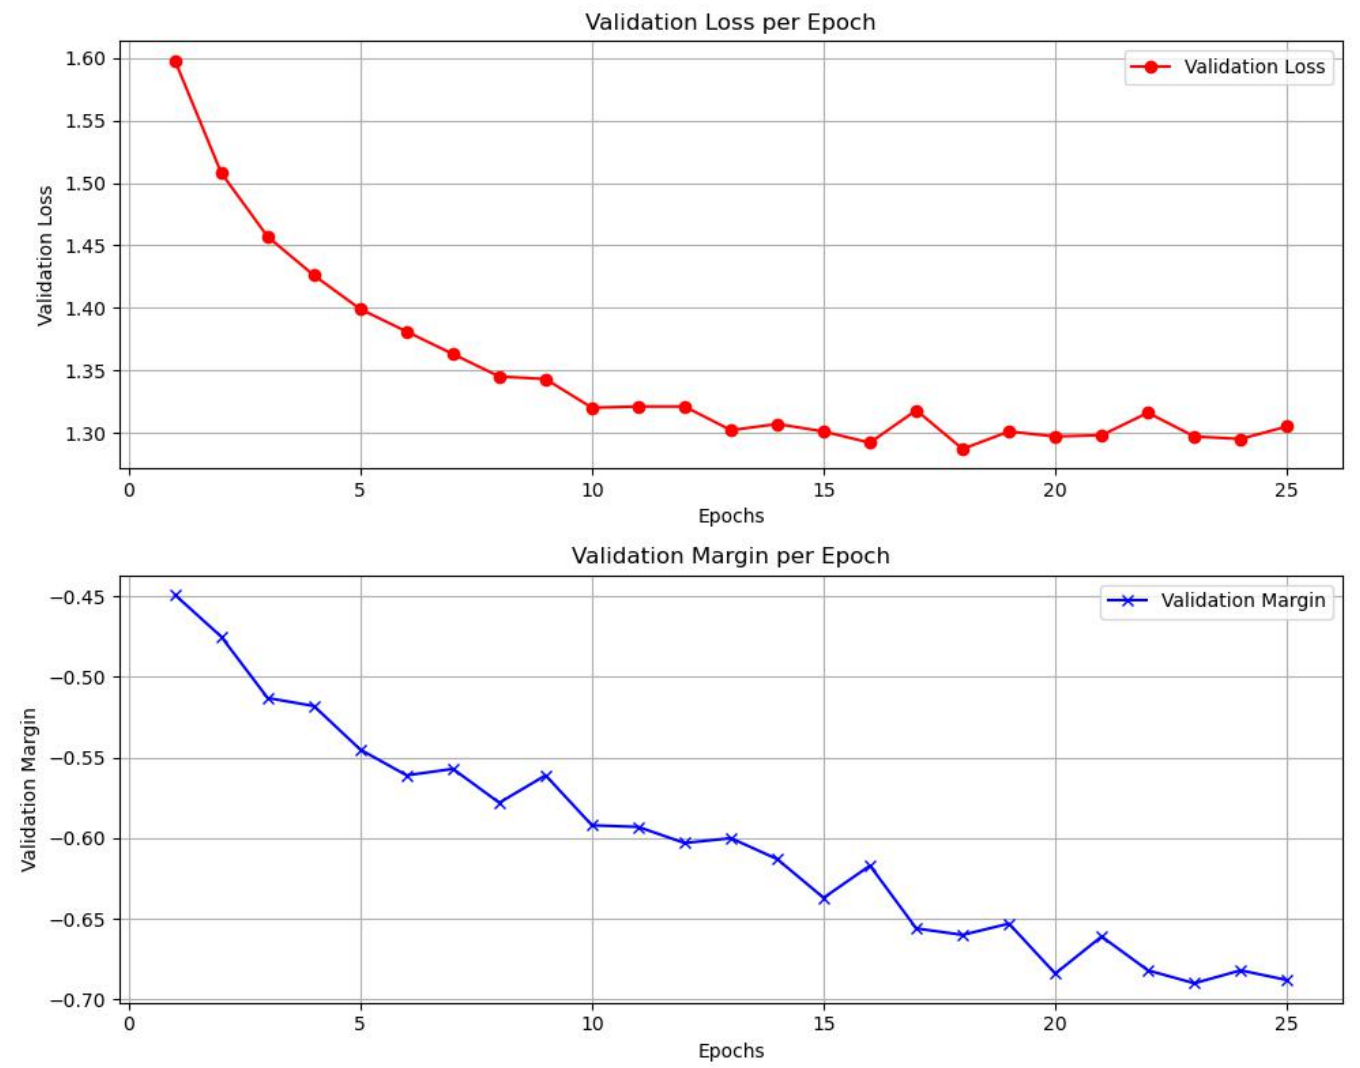


Model2

```
===== Model 2 =====
Training loss: 0.967   Training margin -0.539   Training accuracy: 0.672
Validation accuracy: 0.554

Frobenius1:      12.3
Frobenius2:      5.49
Distance1:       8.15
Distance2:       4.2
Spectral1:       1.87
Spectral2:       2.76
Fro_Fro:        67.4
L1max_L1max:    1.25e+03
Spec_Dist:      24.9
Dist_Spec:      3.59e+02
Spec_Dist_sum:  3.84e+02
Spec_L1max:     24.7
L1max_Spec:     3.53e+02
Spec_L1max_sum: 3.78e+02
Dist_Fro:       44.7
#parameter:     7.89e+05
```

VC bound:	2.03e+09
L1max bound:	1.2e+11
Your bound:	3.84e+10



In general, `model1` has more hidden units which make the model to learn more information from the input hence why the margin and bound greater than `model2`. If we want to minimize the the difference between the two models, we can use regularization to increase the margin of `model2`.